

## Reminder: Q-learning

For transition  $(s, a, r, s')$ :

$$Q_{k+1}^*(s, a) \leftarrow Q_k^*(s, a) + \alpha_k \underbrace{\left( r + \gamma \max_{a'} Q_k^*(s', a') - Q_k^*(s, a) \right)}_{\text{temporal difference}}$$

**Bellman target**

## Reminder: Q-learning

For transition  $(s, a, r, s')$ :

$$Q_{k+1}^*(s, a) \leftarrow Q_k^*(s, a) + \alpha_k \underbrace{\left( r + \gamma \max_{a'} Q_k^*(s', a') - Q_k^*(s, a) \right)}_{\text{temporal difference}}$$

**Bellman target**

**temporal difference**

**Q-learning** can learn  $Q^*$  from trial and error in *finite* MDPs:  $|\mathcal{S}| \ll \infty, |\mathcal{A}| \ll \infty$

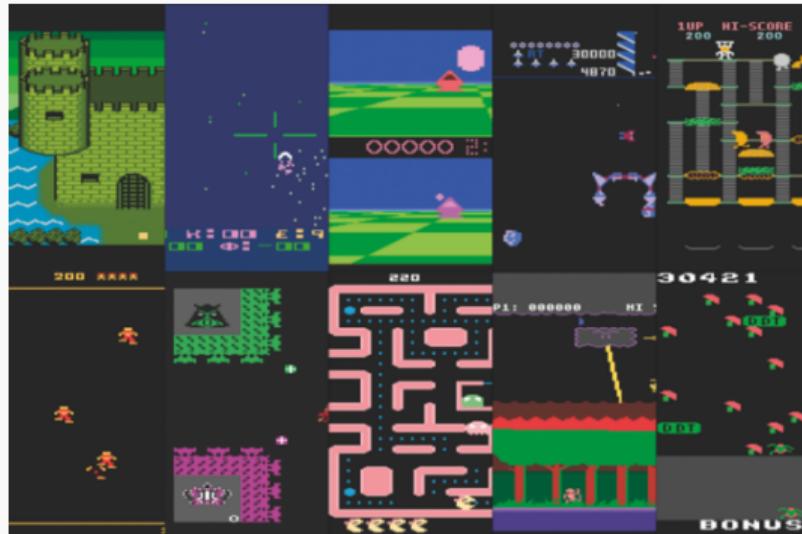
**Important!**

**Q-learning is off-policy**

# Atari Games

Setup: ~~TS  $\ll \infty$~~ ,  $|\mathcal{A}| \ll \infty$

- 57 various games
- Only screen image as input.
- No game-specific features.
- Finite-state case... not quite finite.
- $|\mathcal{A}| \leq 18$



# Atari Games

Setup: ~~TS  $\ll \infty$~~ ,  $|\mathcal{A}| \ll \infty$

- 57 various games
- Only screen image as input.
- No game-specific features.
- Finite-state case... not quite finite.
- $|\mathcal{A}| \leq 18$



Approximate  $Q^*(s, a)$  with neural network!

# Is Atari game a MDP?



# Practical notes: preprocessing

## Action selection frequency:

- Framestack;
- Frameskip;
- MaxAndSkip;
- (sometimes) Sticky actions;

## Atari-specific preprocessing:

- EpisodicLife;
- FireReset;

## Standard tricks:

- Crop image;
- Rescale (often to 84x84);
- Grayscale;

## Reward preprocessing:

- (!) Clip reward to  $\{-1, 0, 1\}$ ;

# Practical notes: preprocessing

## Action selection frequency:

- Framestack;
- Frameskip;
- MaxAndSkip;
- (sometimes) Sticky actions;

## Atari-specific preprocessing:

- EpisodicLife;
- FireReset;

## Standard tricks:

- Crop image;
- Rescale (often to 84x84);
- Grayscale;

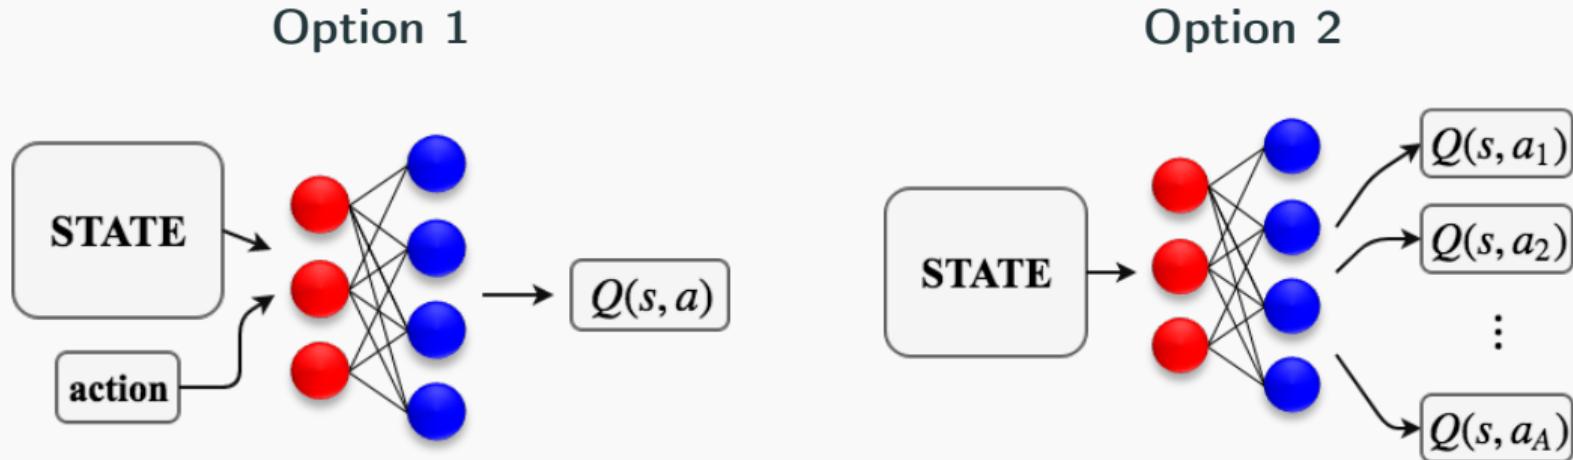
## Reward preprocessing:

- (!) Clip reward to  $\{-1, 0, 1\}$ ;

Important!

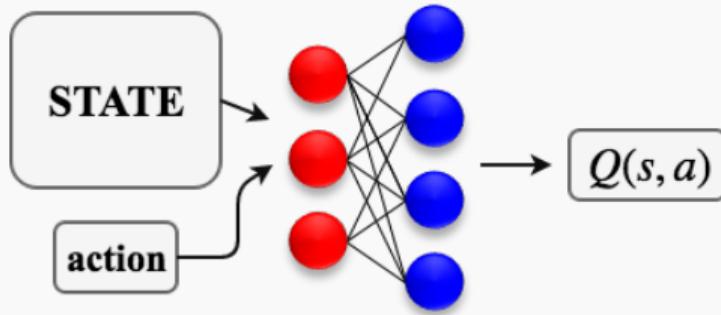
No hyperparameter tuning for each specific game!

# Deep Q-network

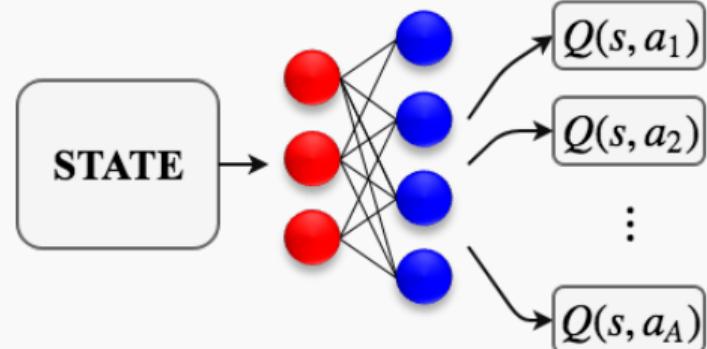


# Deep Q-network

Option 1



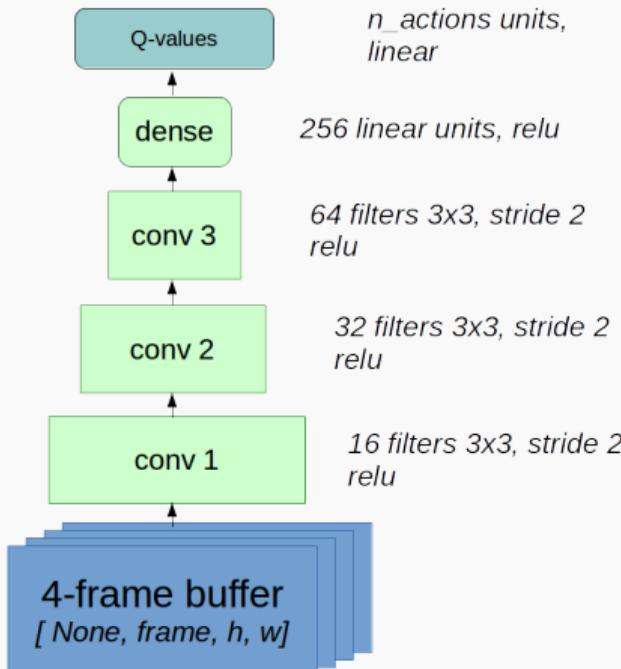
Option 2



✗  $\underset{a}{\operatorname{argmax}} Q^*(s, a, \theta)$  — expensive!

✓  $\underset{a}{\operatorname{argmax}} Q^*(s, a, \theta)$  — one forward pass.

# Deep Q-network: architecture



- 3-4 convolutional layers followed by 1-2 dense layers;
- stride  $> 1$  for size reduction; linear layers still wide;

**Think twice before using:**

- ✗ max pooling
- ✗ batch normalization
- ✗ dropout

## Idea: Approximate Dynamic Programming

Consider the following **regression task**:

- $s, a$  is input;
- $y \in \mathbb{R}$  is target:

$$y(s, a) := r(s, a) + \gamma \mathbb{E}_{s'} \max_{a'} Q^*(s', a', \theta_k)$$

## Idea: Approximate Dynamic Programming

Consider the following **regression task**:

- $s, a$  is input;
- $y \in \mathbb{R}$  is target:

$$y(s, a) := r(s, a) + \gamma \mathbb{E}_{s'} \max_{a'} Q^*(s', a', \theta_k)$$

- MSE loss function:  $\text{Loss}(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$

## Idea: Approximate Dynamic Programming

Consider the following **regression task**:

- $s, a$  is input;
- $y \in \mathbb{R}$  is target:

$$y(s, a) := r(s, a) + \gamma \mathbb{E}_{s'} \max_{a'} Q^*(s', a', \theta_k)$$

- MSE loss function:  $\text{Loss}(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$



$$\frac{1}{2} \mathbb{E}_{(s, a, y)} (y - Q^*(s, a, \theta_{k+1}))^2 \rightarrow \min_{\theta_{k+1}}$$

## Looking at the gradient

$$\nabla_{\theta} \frac{1}{2} (y - Q^*(s, a, \theta))^2 =$$

## Looking at the gradient

$$\nabla_{\theta} \frac{1}{2} (y - Q^*(s, a, \theta))^2 = \underbrace{(y - Q^*(s, a, \theta))}_{\text{temporal difference}} \overbrace{\nabla_{\theta} Q^*(s, a, \theta)}^{\text{how to increase output value}} =$$

=

## Looking at the gradient

$$\nabla_{\theta} \frac{1}{2} (y - Q^*(s, a, \theta))^2 = \underbrace{(y - Q^*(s, a, \theta))}_{\text{temporal difference}} \overbrace{\nabla_{\theta} Q^*(s, a, \theta)}^{\text{how to increase output value}} =$$
$$= \left( r + \gamma \mathbb{E}_{s'} \max_{a'} Q^*(s', a', \theta_k) - Q^*(s, a, \theta) \right) \nabla_{\theta} Q^*(s, a, \theta)$$

## Looking at the gradient

$$\begin{aligned} \nabla_{\theta} \frac{1}{2} (y - Q^*(s, a, \theta))^2 &= \underbrace{(y - Q^*(s, a, \theta))}_{\text{temporal difference}} \overbrace{\nabla_{\theta} Q^*(s, a, \theta)}^{\text{how to increase output value}} = \\ &= \left( r + \gamma \mathbb{E}_{s'} \max_{a'} Q^*(s', a', \theta_k) - Q^*(s, a, \theta) \right) \nabla_{\theta} Q^*(s, a, \theta) \approx \\ &\approx \left( r + \gamma \max_{a'} Q^*(s', a', \theta_k) - Q^*(s, a, \theta) \right) \nabla_{\theta} Q^*(s, a, \theta), \end{aligned}$$

where  $s' \sim p(s' | s, a)$ .

## Deep Q-learning

Consider the following **regression task**:

- $s, a$  is input;
- $y \in \mathbb{R}$  is target:

$$y(s, a) := r + \gamma \max_{a'} Q^*(s', a', \theta_k)$$

where  $r = r(s, a), s' \sim p(s' | s, a);$

## Deep Q-learning

Consider the following **regression task**:

- $s, a$  is input;
- $y \in \mathbb{R}$  is target:

$$y(s, a) := r + \gamma \max_{a'} Q^*(s', a', \theta_k)$$

where  $r = r(s, a), s' \sim p(s' | s, a);$

- MSE loss function:  $\text{Loss}(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$

## Target network

For how long to solve one regression task?

## Target network

For how long to solve one regression task?

- One SGD iteration

## Target network

For how long to solve one regression task?

- One SGD iteration
  - ✗ completely unstable :(

## Target network

For how long to solve one regression task?

- One SGD iteration
  - × completely unstable :(
- Till convergence

## Target network

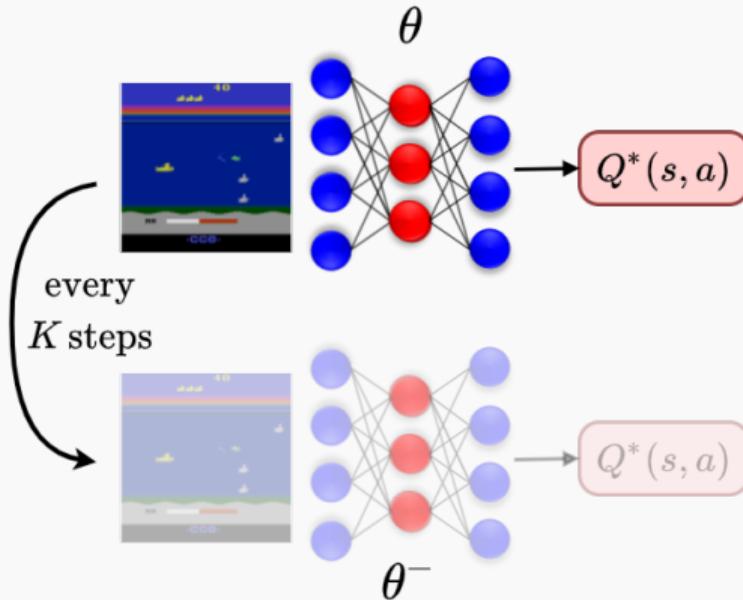
For how long to solve one regression task?

- One SGD iteration
  - ✗ completely unstable :(
- Till convergence
  - ✗ too long
- ✓  $\approx$ 1000 SGD iterations

# Target network

For how long to solve one regression task?

- One SGD iteration
  - ✗ completely unstable :(
- Till convergence
  - ✗ too long
- ✓  $\approx 1000$  SGD iterations

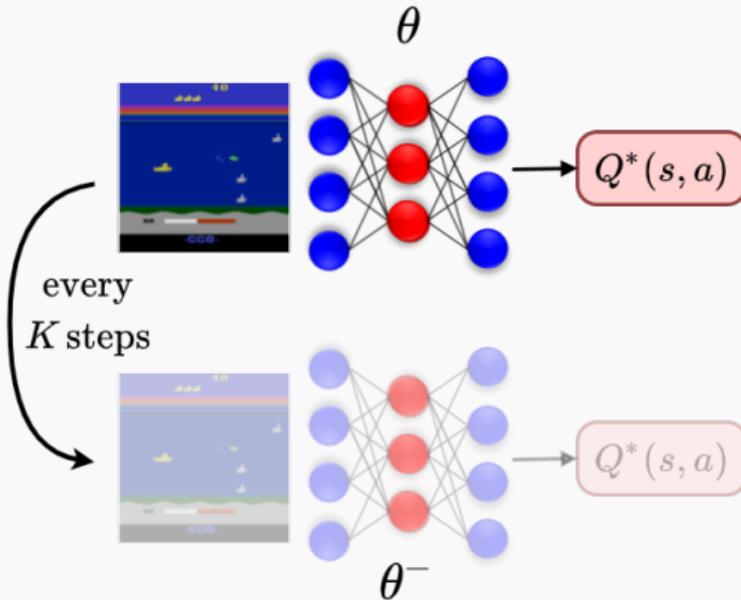


Store a copy of your old Q-network  $Q^*(s, a, \theta^-)$

# Target network

For how long to solve one regression task?

- One SGD iteration
  - ✗ completely unstable :(
- Till convergence
  - ✗ too long
- ✓  $\approx 1000$  SGD iterations
  - $\theta^- \leftarrow \theta$  every  $K$  SGD iterations
  - $\theta^- \leftarrow (1 - \beta)\theta^- + \beta\theta$



Store a copy of your old Q-network  $Q^*(s, a, \theta^-)$

# Sample Decorrelation



# Experience Replay



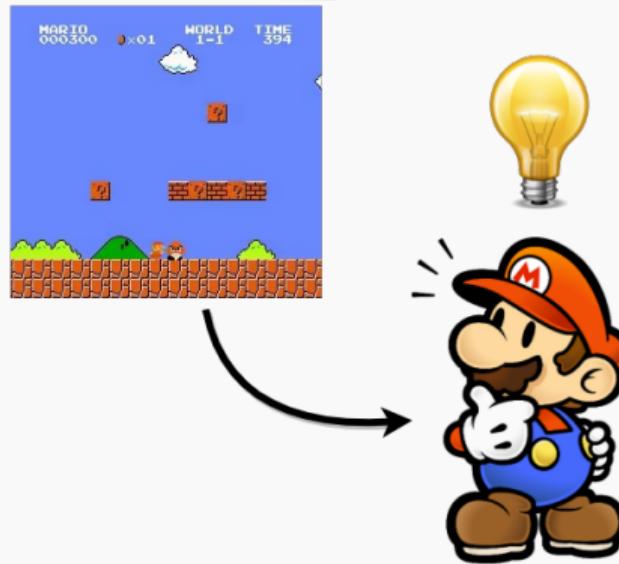
# Do we still need exploration?

## *Trial and error*



# Do we still need exploration?

## Trial and error

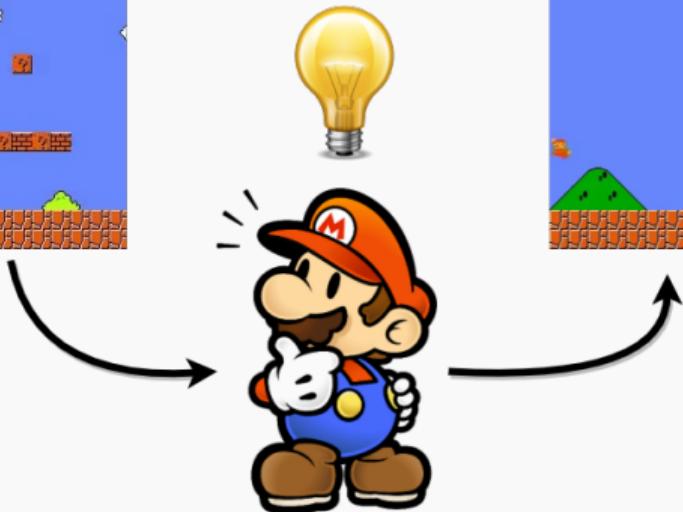


# Do we still need exploration?

*Trial and error*



*Local optima*

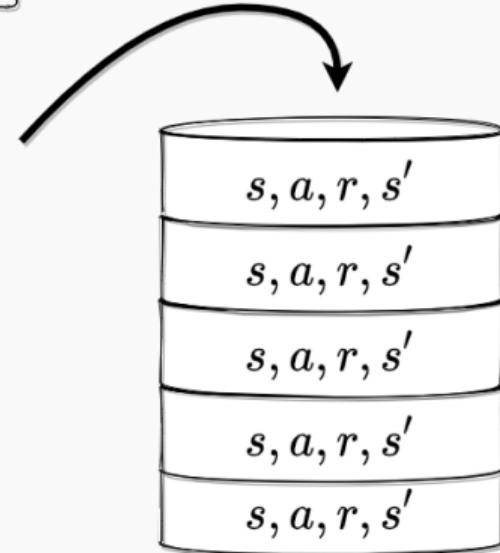


# Do we still need exploration?

*Trial and error*



*Local optima*

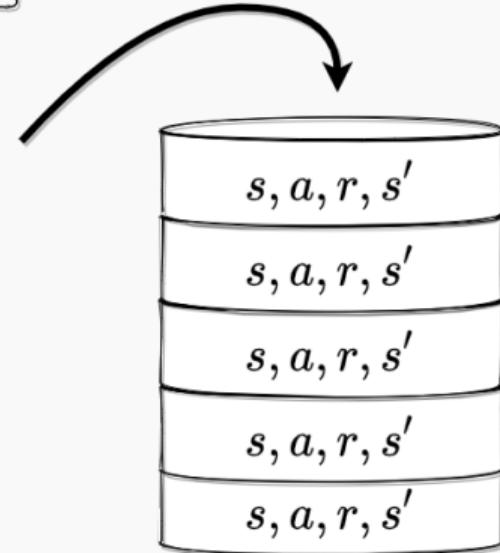
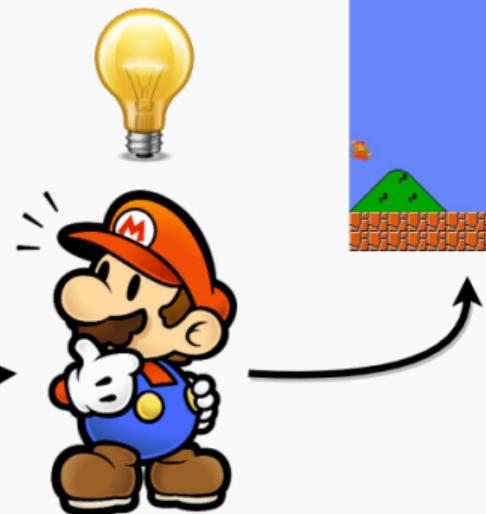


# Do we still need exploration?

## Trial and error



## Local optima



Important !

Diverse data is still required!

# Full alrogithm

## Deep Q-learning

**Initialize**  $Q^*(s, a, \theta)$  arbitrarily,  $\theta^- := \theta$ ,  $\mathcal{D} = \emptyset$ ;

# Full algorithm

## Deep Q-learning

**Initialize**  $Q^*(s, a, \theta)$  arbitrarily,  $\theta^- := \theta$ ,  $\mathcal{D} = \emptyset$ ;

observe  $s_0$ ;

**for**  $k = 0, 1, 2 \dots$

- take action  $a_k \sim \varepsilon\text{-greedy}(Q^*(s_k, a, \theta))$ ;
- observe  $r_k, s_{k+1}, \text{done}_{k+1}$ , store  $(s_k, a_k, r_k, s_{k+1}, \text{done}_{k+1})$  in  $\mathcal{D}$ ;

# Full algorithm

## Deep Q-learning

**Initialize**  $Q^*(s, a, \theta)$  arbitrarily,  $\theta^- := \theta$ ,  $\mathcal{D} = \emptyset$ ;

observe  $s_0$ ;

**for**  $k = 0, 1, 2 \dots$

- take action  $a_k \sim \varepsilon\text{-greedy}(Q^*(s_k, a, \theta))$ ;
- observe  $r_k, s_{k+1}, \text{done}_{k+1}$ , store  $(s_k, a_k, r_k, s_{k+1}, \text{done}_{k+1})$  in  $\mathcal{D}$ ;
- sample batch of transitions  $\mathbb{T} := (s, a, r, s', \text{done})$  from  $\mathcal{D}$ ;

# Full algorithm

## Deep Q-learning

**Initialize**  $Q^*(s, a, \theta)$  arbitrarily,  $\theta^- := \theta$ ,  $\mathcal{D} = \emptyset$ ;

observe  $s_0$ ;

**for**  $k = 0, 1, 2 \dots$

- take action  $a_k \sim \varepsilon\text{-greedy}(Q^*(s_k, a, \theta))$ ;
- observe  $r_k, s_{k+1}, \text{done}_{k+1}$ , store  $(s_k, a_k, r_k, s_{k+1}, \text{done}_{k+1})$  in  $\mathcal{D}$ ;
- sample batch of transitions  $\mathbb{T} := (s, a, r, s', \text{done})$  from  $\mathcal{D}$ ;
- $y(\mathbb{T}) := r + \gamma(1 - \text{done}) \max_{a'} Q^*(s', a', \theta^-)$ ;

# Full algorithm

## Deep Q-learning

**Initialize**  $Q^*(s, a, \theta)$  arbitrarily,  $\theta^- := \theta$ ,  $\mathcal{D} = \emptyset$ ;

observe  $s_0$ ;

**for**  $k = 0, 1, 2 \dots$

- take action  $a_k \sim \varepsilon\text{-greedy}(Q^*(s_k, a, \theta))$ ;
- observe  $r_k, s_{k+1}, \text{done}_{k+1}$ , store  $(s_k, a_k, r_k, s_{k+1}, \text{done}_{k+1})$  in  $\mathcal{D}$ ;
- sample batch of transitions  $\mathbb{T} := (s, a, r, s', \text{done})$  from  $\mathcal{D}$ ;
- $y(\mathbb{T}) := r + \gamma(1 - \text{done}) \max_{a'} Q^*(s', a', \theta^-)$ ;
- perform a step of gradient descent:

$$\theta \leftarrow \theta - \frac{\alpha}{B} \sum_{\mathbb{T}} \nabla_{\theta} (Q^*(s, a, \theta) - y(\mathbb{T}))^2$$

# Full algorithm

## Deep Q-learning

**Initialize**  $Q^*(s, a, \theta)$  arbitrarily,  $\theta^- := \theta$ ,  $\mathcal{D} = \emptyset$ ;

observe  $s_0$ ;

**for**  $k = 0, 1, 2 \dots$

- take action  $a_k \sim \varepsilon\text{-greedy}(Q^*(s_k, a, \theta))$ ;
- observe  $r_k, s_{k+1}, \text{done}_{k+1}$ , store  $(s_k, a_k, r_k, s_{k+1}, \text{done}_{k+1})$  in  $\mathcal{D}$ ;
- sample batch of transitions  $\mathbb{T} := (s, a, r, s', \text{done})$  from  $\mathcal{D}$ ;
- $y(\mathbb{T}) := r + \gamma(1 - \text{done}) \max_{a'} Q^*(s', a', \theta^-)$ ;
- perform a step of gradient descent:

$$\theta \leftarrow \theta - \frac{\alpha}{B} \sum_{\mathbb{T}} \nabla_{\theta} (Q^*(s, a, \theta) - y(\mathbb{T}))^2$$

- update target network: if  $k \bmod K = 0$ :  $\theta^- \leftarrow \theta$

There is a lot to improve



Multistep  
DQN



Double  
DQN

DQN



Prioritized  
Replay



Noisy  
Net



Categorical  
DQN



Dueling  
DQN



# Overestimation Bias and how to fight it

**Problem:** model often *overestimates* future reward.

# Overestimation Bias and how to fight it

**Problem:** model often *overestimates* future reward.

$$\max_{a'}(Q^*(s', a', \theta) + \underbrace{\varepsilon(s', a')}_\text{approximation error or luck})$$

# Overestimation Bias and how to fight it

**Problem:** model often *overestimates* future reward.

$$\mathbb{E}_\varepsilon \max_{a'} (Q^*(s', a', \theta) + \underbrace{\varepsilon(s', a')}_\text{approximation error or luck}) \geq \max_{a'} Q^*(s', a')$$

# Overestimation Bias and how to fight it

**Problem:** model often *overestimates* future reward.

$$\mathbb{E}_\varepsilon \max_{a'} (Q^*(s', a', \theta) + \underbrace{\varepsilon(s', a')}_\text{approximation error or luck}) \geq \max_{a'} Q^*(s', a')$$

Decouple **action selection** and **action evaluation**:



$$\max_{a'} Q^*(s', a') = \overbrace{Q^*(s', \underbrace{\operatorname{argmax}_{a'} Q^*(s', a')})}^\text{action evaluation}$$

action selection

# Overestimation Bias and how to fight it

**Problem:** model often *overestimates* future reward.

$$\mathbb{E}_\varepsilon \max_{a'} (Q^*(s', a', \theta) + \underbrace{\varepsilon(s', a')}_\text{approximation error or luck}) \geq \max_{a'} Q^*(s', a')$$

Decouple **action selection** and **action evaluation**:



$$\max_{a'} Q^*(s', a') = \overbrace{Q^*(s', \underset{a'}{\operatorname{argmax}} Q^*(s', a'))}^{\text{action evaluation}} \approx Q_1^*(s', \underset{a'}{\operatorname{argmax}} Q_2^*(s', a'))$$

action selection

## Action Evaluation: options

Name	Networks	Targets
Double Q-learning not really	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \max_{a'} Q_?^*(s', a')$

## Action Evaluation: options

Name	Networks	Targets
Double Q-learning not really	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \max_{a'} Q_2^*(s', a')$

## Action Evaluation: options

Name	Networks	Targets
Double Q-learning not really	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \max_{a'} Q_2^*(s', a')$ $y_2 = r + \gamma \max_{a'} Q_1^*(s', a')$

## Action Evaluation: options

Name	Networks	Targets
Double Q-learning Twin DQN still not really	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \max_{a'} Q_2^*(s', a')$ $y_2 = r + \gamma \max_{a'} Q_1^*(s', a')$
Double DQN	$Q^*$ $Q_-^*$ — target network	

## Action Evaluation: options

Name	Networks	Targets
Double Q-learning Twin DQN still not ready	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \max_{a'} Q_2^*(s', a')$ $y_2 = r + \gamma \max_{a'} Q_1^*(s', a')$
Double DQN	$Q^*$ $Q_-^*$ — target network	$y = r + \gamma \max_{a'} Q_-^*(s', a')$

## Action Evaluation: options

Name	Networks	Targets
Double Q-learning Twin DQN ?!?	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \max_{a'} Q_2^*(s', a')$ $y_2 = r + \gamma \max_{a'} Q_1^*(s', a')$
Double DQN	$Q^*$ $Q_-^*$ — target network	$y = r + \gamma \max_{a'} Q_-^*(s', a')$
Twin DQN («Clipped Double DQN»)	$Q_1^*$ $Q_2^*$	

## Action Evaluation: options

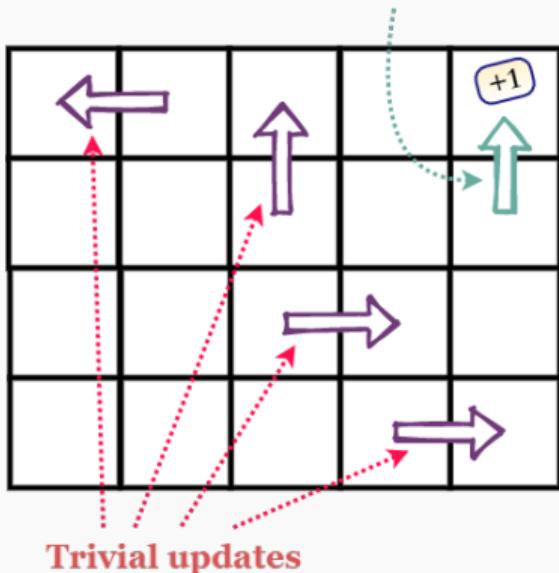
Name	Networks	Targets
Double Q-learning Twin DQN ?!?	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \max_{a'} Q_2^*(s', a')$ $y_2 = r + \gamma \max_{a'} Q_1^*(s', a')$
Double DQN	$Q^*$ $Q_-^*$ — target network	$y = r + \gamma \max_{a'} Q_-^*(s', a')$
Twin DQN («Clipped Double DQN»)	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \min_{i=1,2} Q_i^*(s', \arg\max_{a'} Q_1^*(s', a'))$

## Action Evaluation: options

Name	Networks	Targets
<del>Double Q learning</del> <del>Twin DQN</del> ?!?	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \max_{a'} Q_2^*(s', a')$ $y_2 = r + \gamma \max_{a'} Q_1^*(s', a')$
Double DQN	$Q^*$ $Q_-^*$ — target network	$y = r + \gamma \max_{a'} Q_-^*(s', a')$
Twin DQN («Clipped Double DQN»)	$Q_1^*$ $Q_2^*$	$y_1 = r + \gamma \min_{i=1,2} Q_i^*(s', \arg\max_{a'} Q_1^*(s', a'))$ $y_2 = r + \gamma \min_{i=1,2} Q_i^*(s', \arg\max_{a'} Q_2^*(s', a'))$

# Prioritized Experience Replay

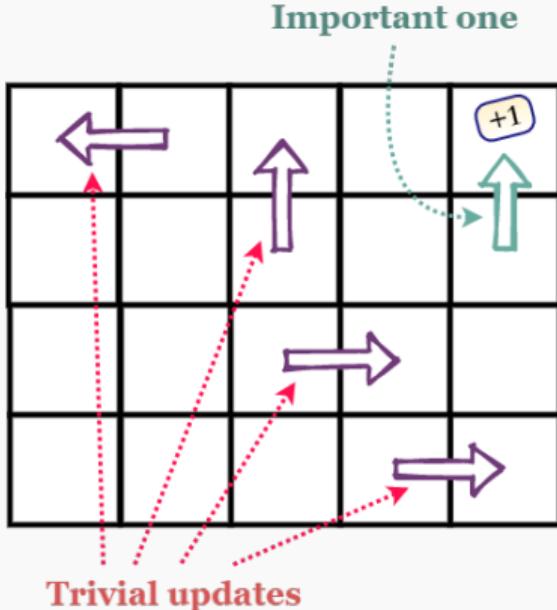
Important one



**Problem:** a lot of trivial updates.

**Goal:** propagate reinforcement from future to past

# Prioritized Experience Replay



**Problem:** a lot of trivial updates.

**Goal:** propagate reinforcement from future to past

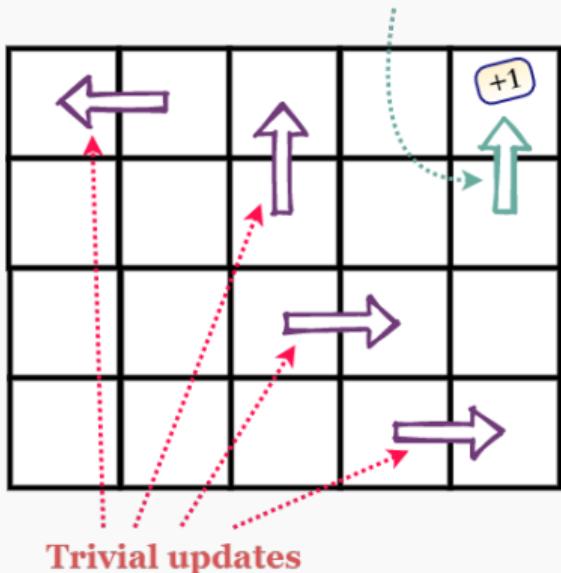


Prioritized sampling from replay buffer:

$$P(T) \propto |y(T) - Q^*(s, a, \theta)|^\alpha$$

# Prioritized Experience Replay

Important one



**Problem:** a lot of trivial updates.

**Goal:** propagate reinforcement from future to past

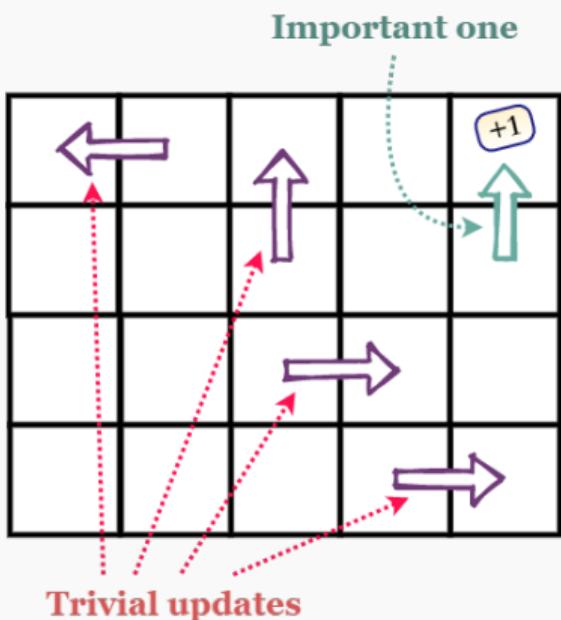


Prioritized sampling from replay buffer:

$$P(T) \propto |y(T) - Q^*(s, a, \theta)|^\alpha$$

Issues?

# Prioritized Experience Replay



**Problem:** a lot of trivial updates.

**Goal:** propagate reinforcement from future to past



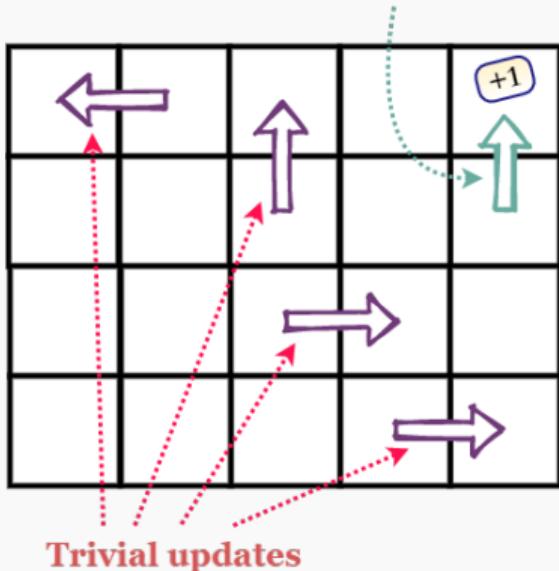
Prioritized sampling from replay buffer:

$$P(T) \propto |y(T) - Q^*(s, a, \theta)|^\alpha$$

**Issues?**  $s' \not\sim p(s' | s, a)$  now.

# Prioritized Experience Replay

Important one



**Problem:** a lot of trivial updates.

**Goal:** propagate reinforcement from future to past



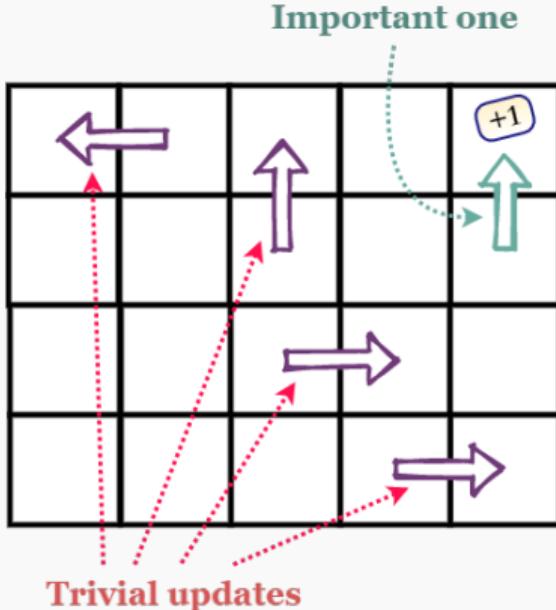
Prioritized sampling from replay buffer:

$$P(T) \propto |y(T) - Q^*(s, a, \theta)|^\alpha$$

**Issues?**  $s' \not\sim p(s' | s, a)$  now.

$$\mathbb{E}_{T \sim \text{Uniform}} \text{Loss}(T) \propto$$

# Prioritized Experience Replay



**Problem:** a lot of trivial updates.

**Goal:** propagate reinforcement from future to past



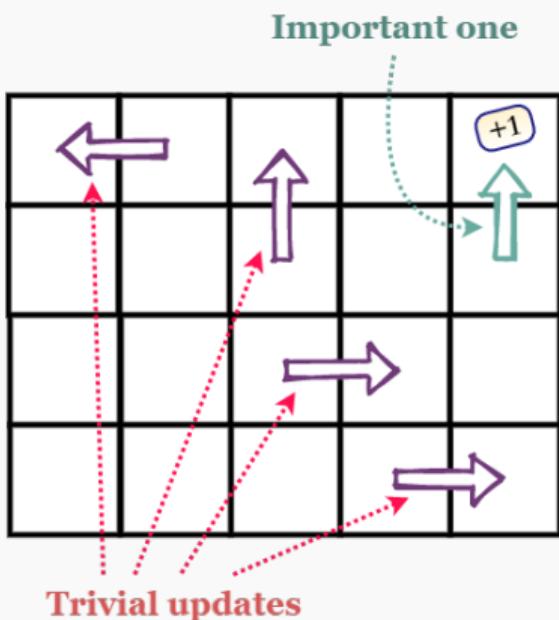
Prioritized sampling from replay buffer:

$$P(T) \propto |y(T) - Q^*(s, a, \theta)|^\alpha$$

**Issues?**  $s' \not\sim p(s' | s, a)$  now.

$$\mathbb{E}_{T \sim \text{Uniform}} \text{Loss}(T) \propto \mathbb{E}_{T \sim P(T)} \underbrace{\frac{1}{P(T)}}_{\text{weight}} \text{Loss}(T),$$

# Prioritized Experience Replay



**Problem:** a lot of trivial updates.

**Goal:** propagate reinforcement from future to past



Prioritized sampling from replay buffer:

$$P(\mathbb{T}) \propto |y(\mathbb{T}) - Q^*(s, a, \theta)|^\alpha$$

**Issues?**  $s' \not\sim p(s' | s, a)$  now.

$$\mathbb{E}_{\mathbb{T} \sim \text{Uniform}} \text{Loss}(\mathbb{T}) \approx \mathbb{E}_{\mathbb{T} \sim P(\mathbb{T})} \left( \frac{1}{P(\mathbb{T})} \right)^{\beta(t)} \text{Loss}(\mathbb{T}),$$

where  $\beta(t)$  anneals from 0 to 1.

## Technical note: SumTree structure

Use **SumTree** structure:

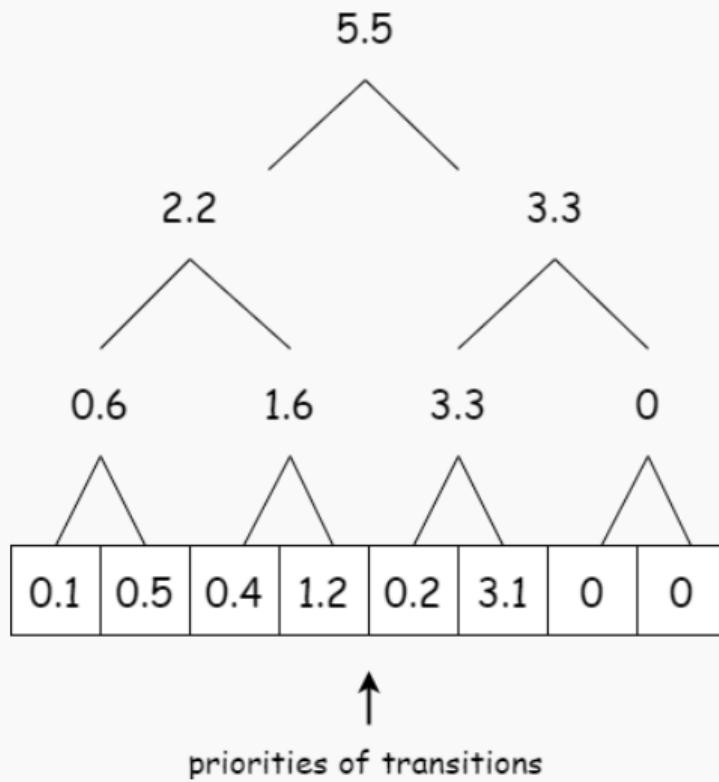
- leaf nodes are sampling priorities

0.1	0.5	0.4	1.2	0.2	3.1	0	0
-----	-----	-----	-----	-----	-----	---	---



priorities of transitions

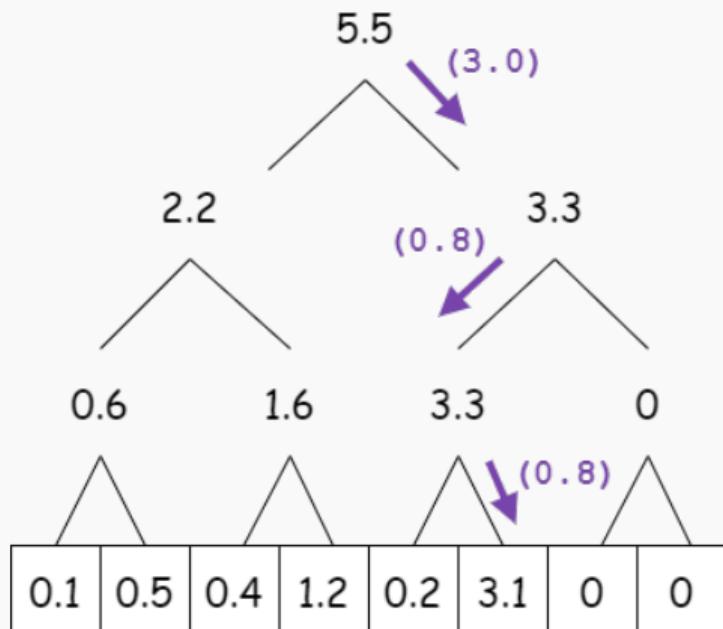
## Technical note: SumTree structure



Use **SumTree** structure:

- leaf nodes are sampling priorities
- each node is sum of its two children

## Technical note: SumTree structure



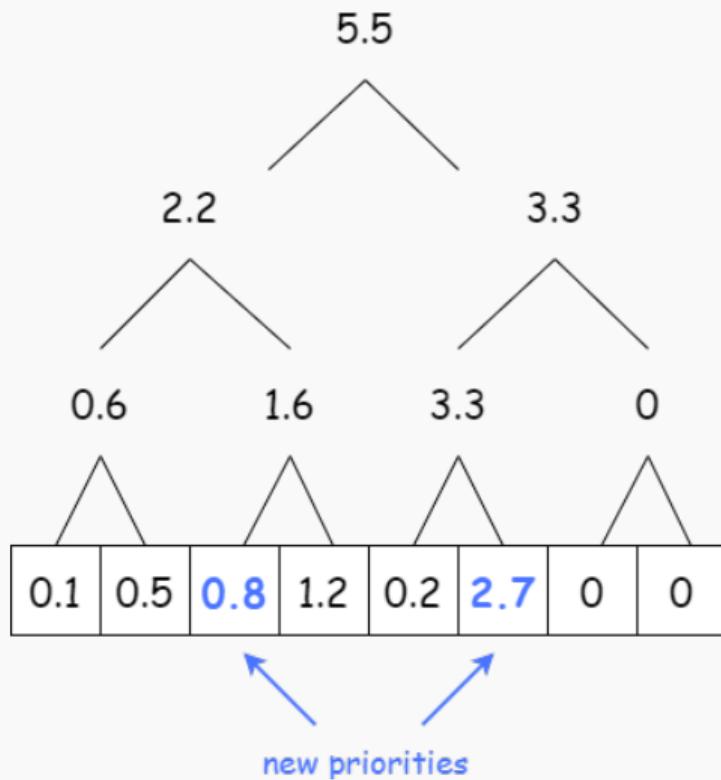
Use **SumTree** structure:

- leaf nodes are sampling priorities
- each node is sum of its two children

**Prioritized buffer procedure:**

- to sample from  $P(\mathbb{T})$  use uniform real number from  $[0, \sum P(\mathbb{T})]$ ;

## Technical note: SumTree structure



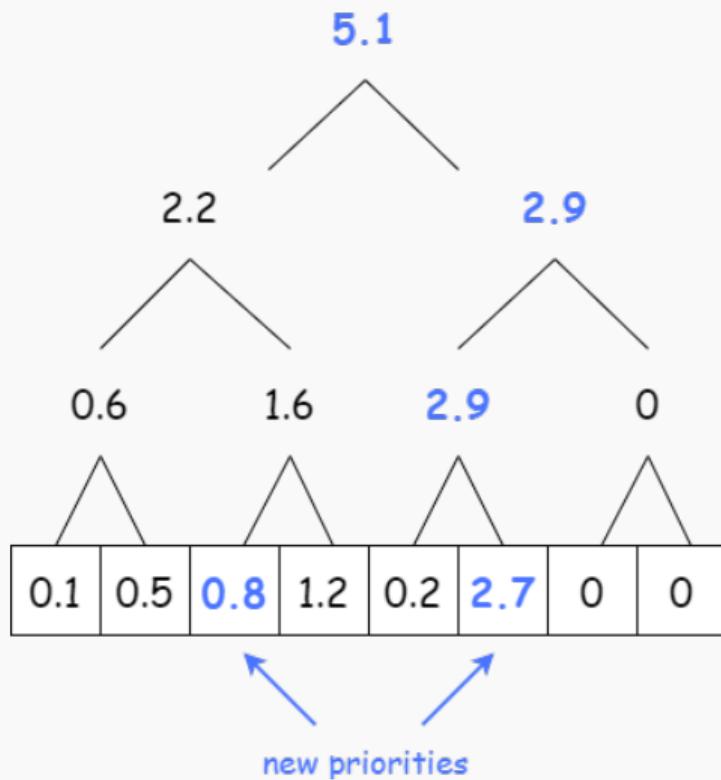
Use **SumTree** structure:

- leaf nodes are sampling priorities
- each node is sum of its two children

**Prioritized buffer procedure:**

- to sample from  $P(\mathbb{T})$  use uniform real number from  $[0, \sum P(\mathbb{T})]$ ;
- train on the batch;
- recompute priorities **for this batch only**;

## Technical note: SumTree structure



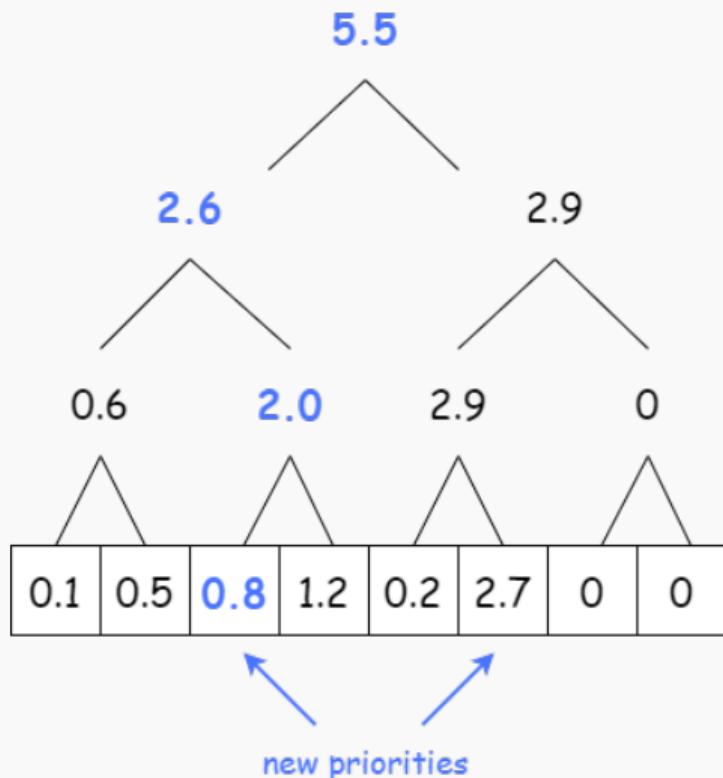
Use **SumTree** structure:

- leaf nodes are sampling priorities
- each node is sum of its two children

**Prioritized buffer procedure:**

- to sample from  $P(\mathbb{T})$  use uniform real number from  $[0, \sum P(\mathbb{T})]$ ;
- train on the batch;
- recompute priorities **for this batch only**;
- update nodes of SumTree;

## Technical note: SumTree structure



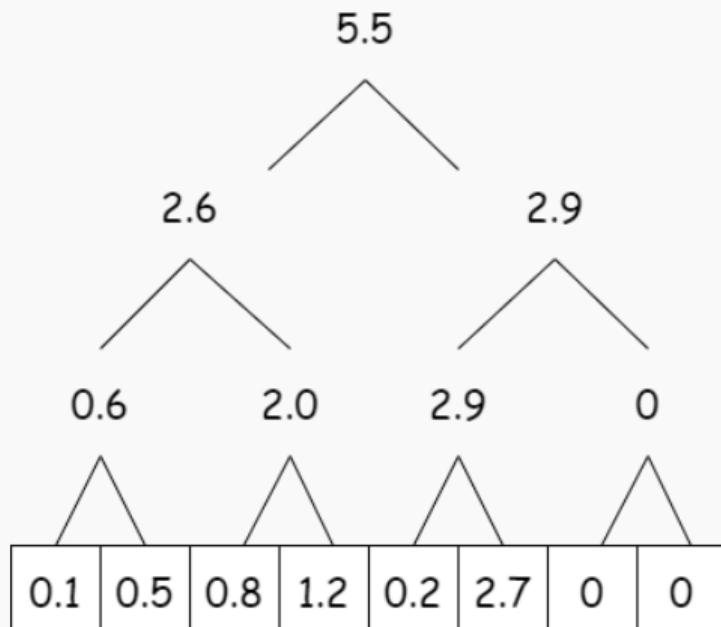
Use **SumTree** structure:

- leaf nodes are sampling priorities
- each node is sum of its two children

**Prioritized buffer procedure:**

- to sample from  $P(\mathbb{T})$  use uniform real number from  $[0, \sum P(\mathbb{T})]$ ;
- train on the batch;
- recompute priorities **for this batch only**;
- update nodes of SumTree;

## Technical note: SumTree structure



Use **SumTree** structure:

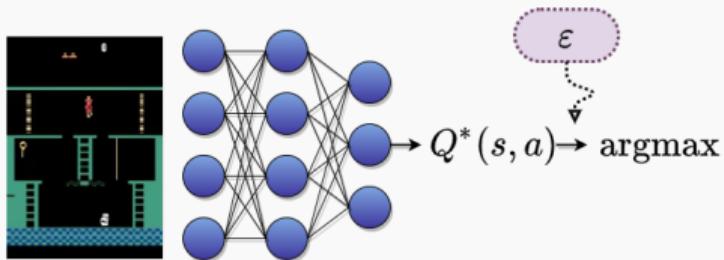
- leaf nodes are sampling priorities
- each node is sum of its two children

**Prioritized buffer procedure:**

- to sample from  $P(\mathbb{T})$  use uniform real number from  $[0, \sum P(\mathbb{T})]$ ;
- train on the batch;
- recompute priorities **for this batch only**;
- update nodes of SumTree;

Complexity:  $O(\log N)$ , where  $N$  is buffer size.

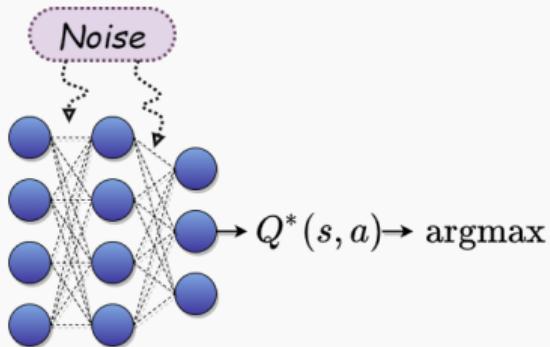
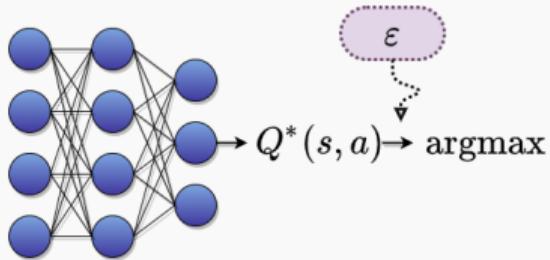
# Noisy Networks



**Problem:**  $\varepsilon$ -greedy is naive

- ✗ inconvenient hyperparameter
- ✗ state-independent exploration

# Noisy Networks



**Problem:**  $\epsilon$ -greedy is naive

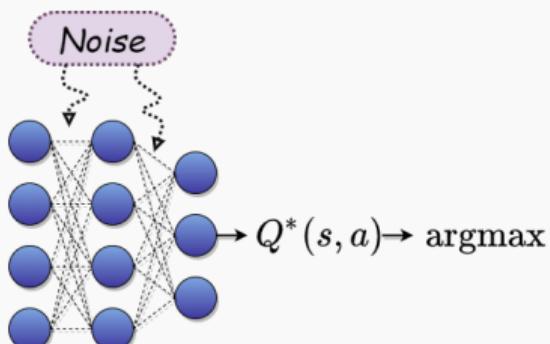
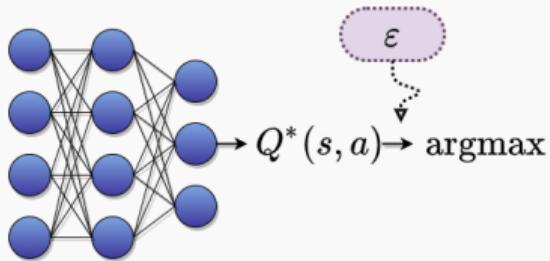
- × inconvenient hyperparameter
- × state-independent exploration

Add noise to the network weights



$$w := \mu + \sigma \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, 1)$$

# Noisy Networks



**Problem:**  $\varepsilon$ -greedy is naive

- ✗ inconvenient hyperparameter
- ✗ state-independent exploration

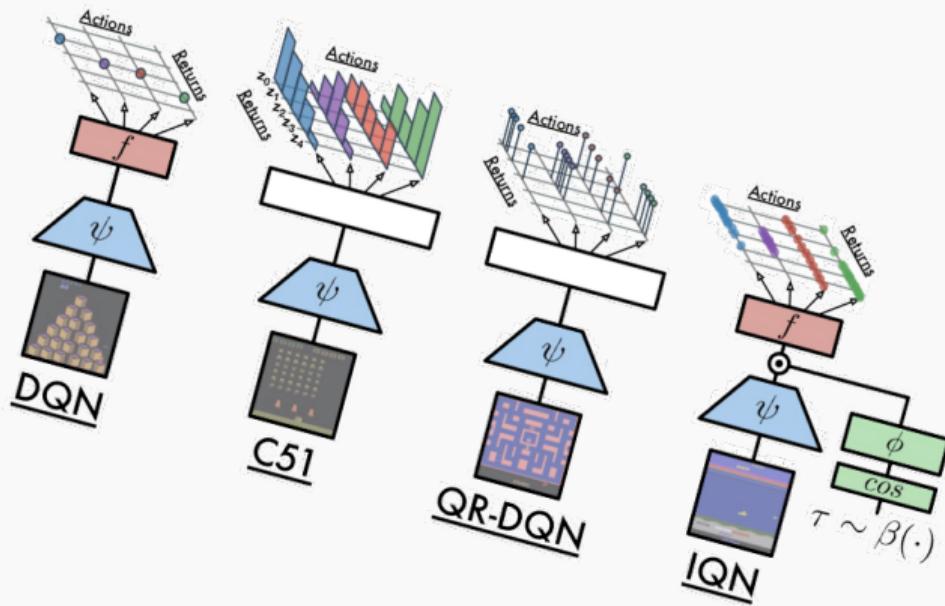
Add noise to the network weights



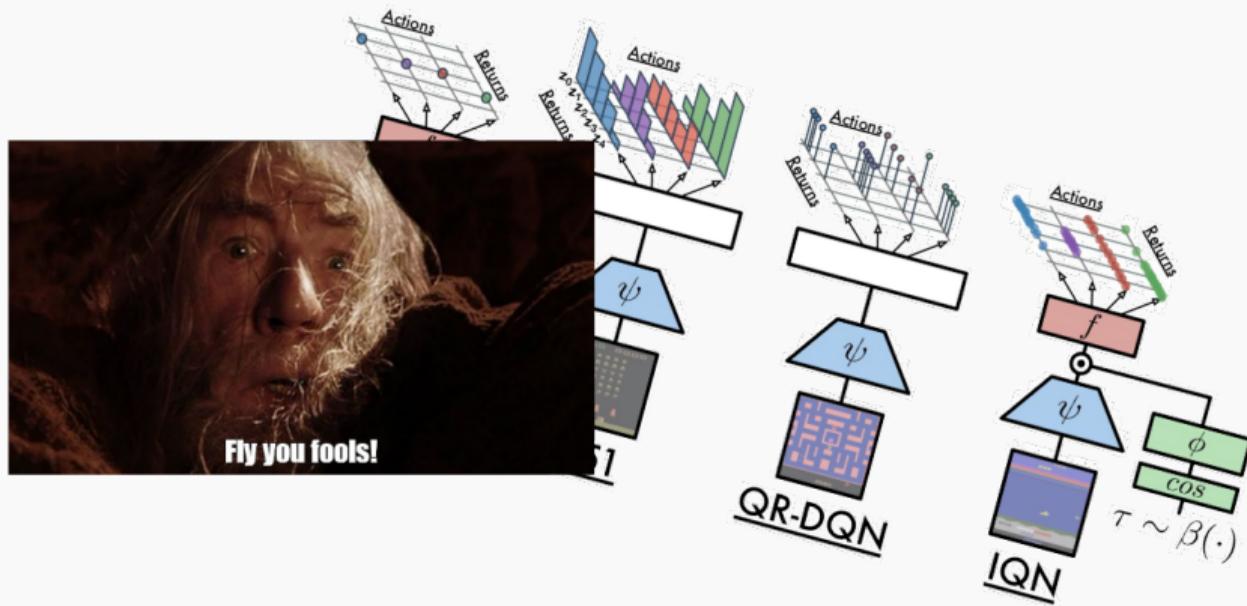
$$w := \mu + \sigma \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, 1)$$

- ✓ no hyperparameters (initialization of  $\sigma$ ?)
- ✓ state-dependent (not interpretable though)
- ✗ expensive in wall-clock time

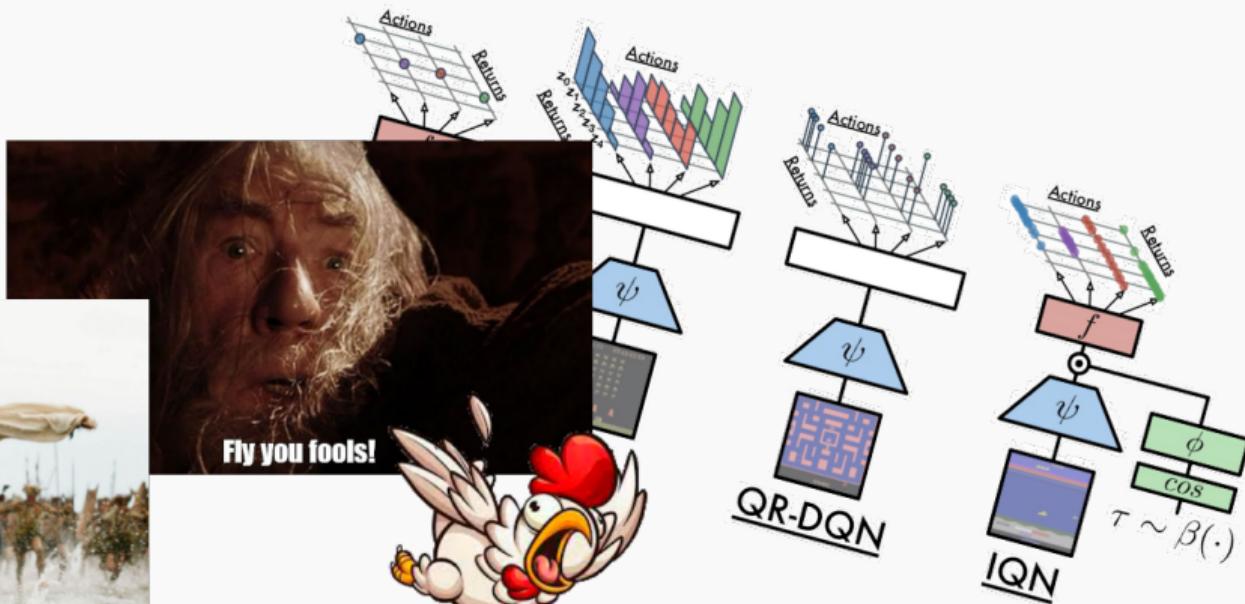
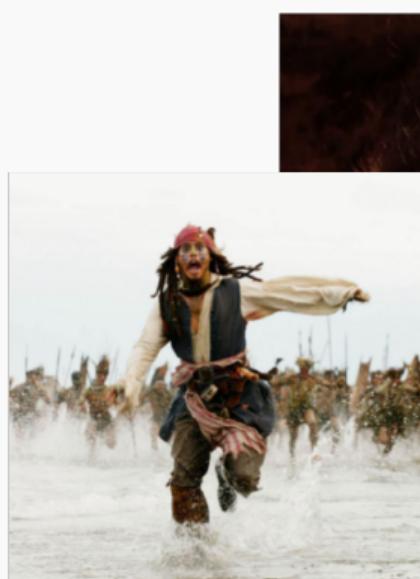
# Distributional RL



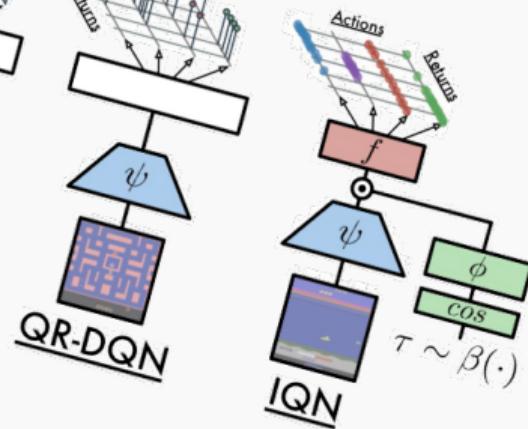
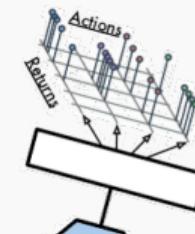
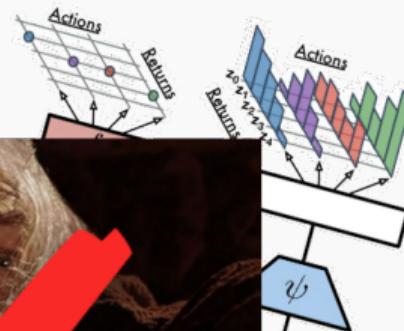
# Distributional RL



# Distributional RL



# Distributional RL



## Dueling DQN

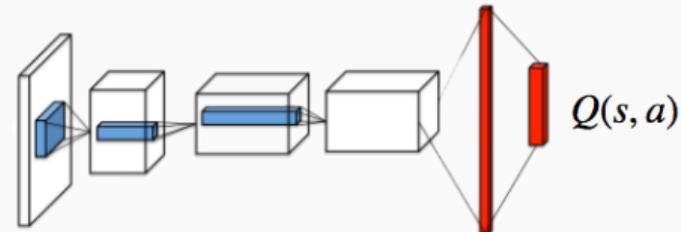
**Problem:** state value estimation requires knowledge about *all* actions

# Dueling DQN

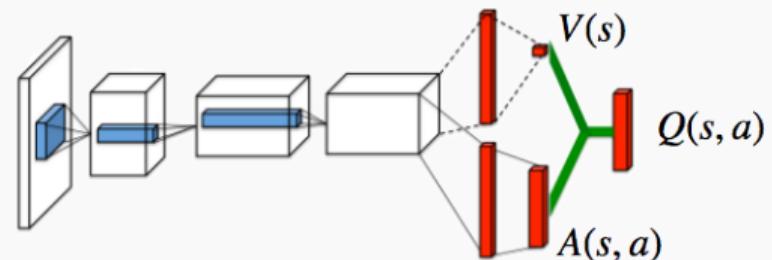
**Problem:** state value estimation requires knowledge about *all* actions

$$Q^*(s, a) := V^*(s) + \underbrace{A^*(s, a)}_{\text{not arbitrary}}$$

**Q-network**



**Dueling Q-network**



# Dueling DQN

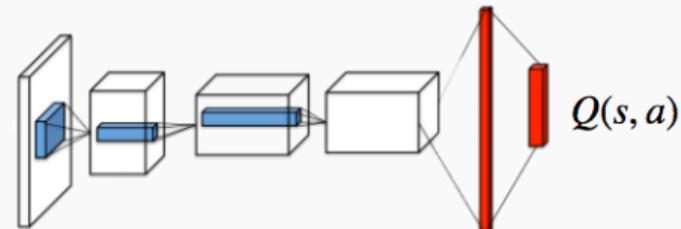
**Problem:** state value estimation requires knowledge about *all* actions

$$Q^*(s, a) := V^*(s) + \underbrace{A^*(s, a)}_{\text{not arbitrary}}$$

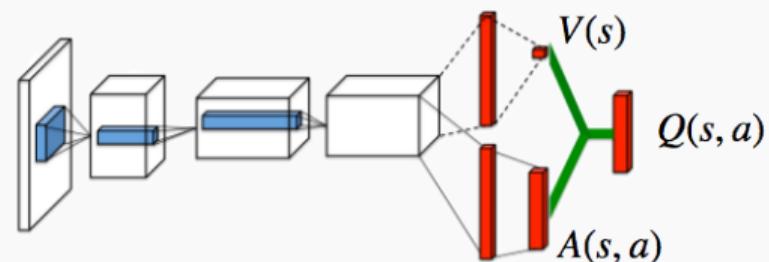
**Danger!**

This introduces one **extra dimension of freedom!**

**Q-network**



**Dueling Q-network**



# Dueling DQN

**Problem:** state value estimation requires knowledge about *all* actions

$$Q^*(s, a) := V^*(s) + \underbrace{A^*(s, a)}_{\text{not arbitrary}}$$

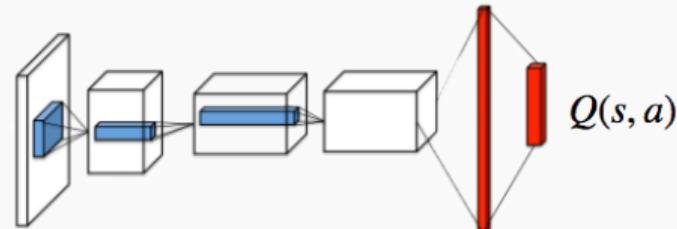
**Danger!**

This introduces one **extra dimension of freedom!**

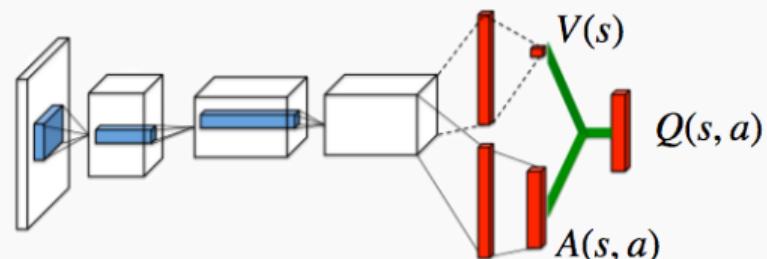
$$Q_\theta^*(s, a) := V_\theta^*(s) + A_\theta^*(s, a) - \underbrace{\text{mean } A_\theta^*(s, a)}_a$$

theory: take max

**Q-network**



**Dueling Q-network**



## Partially Observable MDP

**Problem:** no access to full state description, only to **observations**:

$$o_t \sim p(o_t \mid s_t)$$

## Partially Observable MDP

**Problem:** no access to full state description, only to **observations**:

$$o_t \sim p(o_t | s_t)$$

**Theory:** see **PoMDP** (Partially Observable MDP);

**Practice:** LSTM

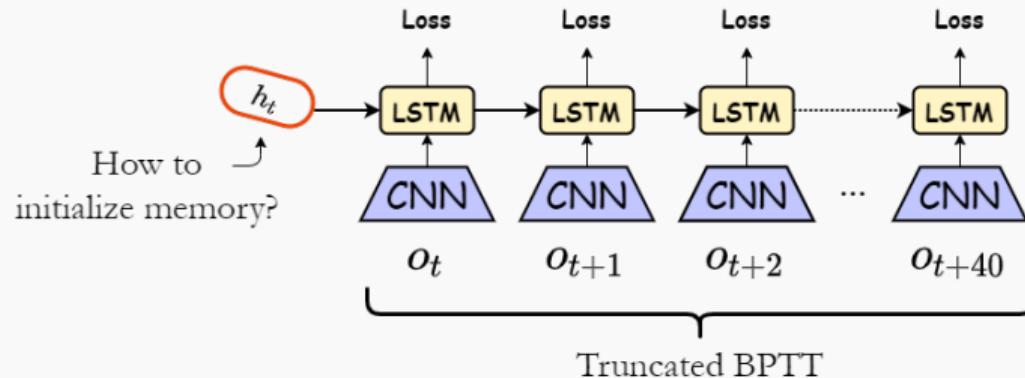
# Partially Observable MDP

**Problem:** no access to full state description, only to **observations**:

$$o_t \sim p(o_t | s_t)$$

**Theory:** see **PoMDP** (Partially Observable MDP);

**Practice:** LSTM



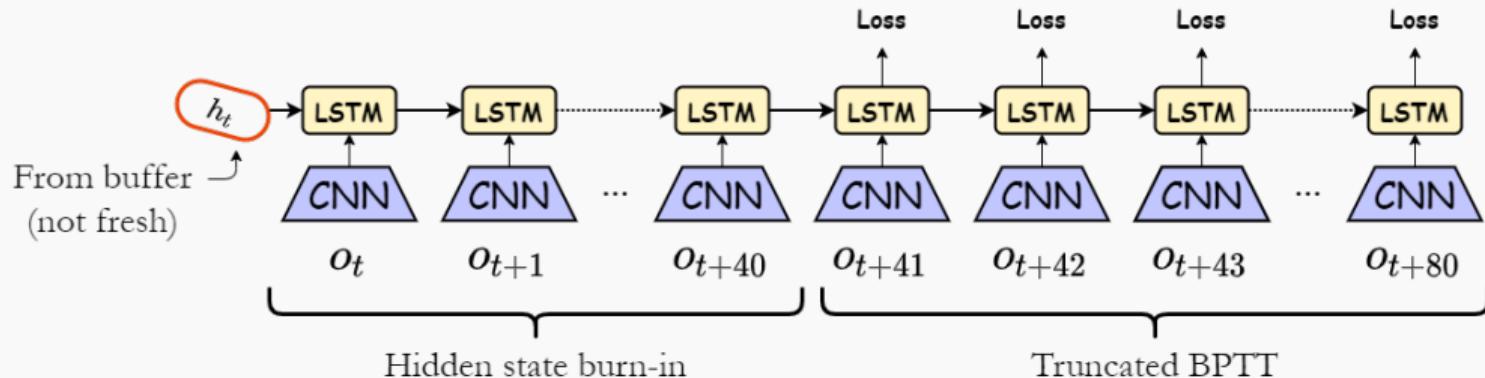
# Partially Observable MDP

**Problem:** no access to full state description, only to **observations**:

$$o_t \sim p(o_t | s_t)$$

**Theory:** see **PoMDP** (Partially Observable MDP);

**Practice:** LSTM



## Multi-step DQN

**Problem:** delayed rewards + compounding error

## Multi-step DQN

**Problem:** delayed rewards + compounding error

$$y := \underbrace{r + \gamma r' + \gamma^2 r'' + \dots}_{\approx \text{«optimal behavior»?}} + \gamma^N \max_{a^{(N)}} Q^*(s^{(N)}, a^{(N)})$$

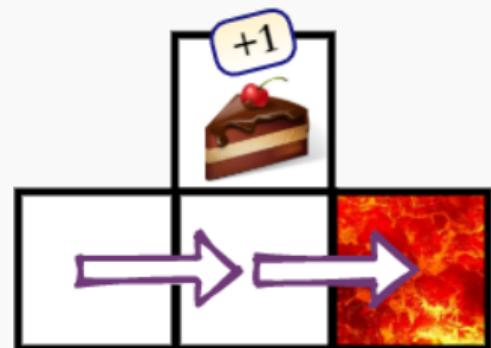
# Multi-step DQN

**Problem:** delayed rewards + compounding error

$$y := \underbrace{r + \gamma r' + \gamma^2 r'' + \dots}_{\approx \text{«optimal behavior»?}} + \gamma^N \max_{a^{(N)}} Q^*(s^{(N)}, a^{(N)})$$

## Important !

- ✓ can *significantly* help with this issue;
- ✗ theoretically **can't be done off-policy**;



# Multi-step DQN

**Problem:** delayed rewards + compounding error

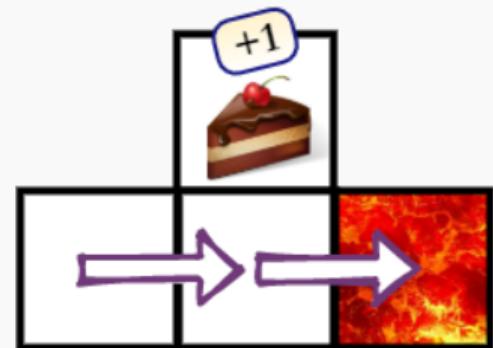
$$y := \underbrace{r + \gamma r' + \gamma^2 r'' + \dots}_{\approx \text{«optimal behavior»?}} + \gamma^N \max_{a^{(N)}} Q^*(s^{(N)}, a^{(N)})$$

## Important!

- ✓ can significantly help with this issue;
- ✗ theoretically can't be done off-policy;

**Implementation:** just store in replay buffer

$$\left( s, \quad a, \quad \sum_{n=0}^{N-1} \gamma^n r^{(n)}, \quad s^{(N)}, \quad \bigvee_{n=1}^N \text{done}_n \right)$$



## Rainbow DQN

(2017)

- Double DQN
- Prioritized buffer
- Multi-step DQN
  - Noisy Nets
  - Dueling DQN
- *Distributional RL*  
*(next time)*

## Rainbow DQN (2017)

- Double DQN
- Prioritized buffer
- Multi-step DQN
  - Noisy Nets
  - Dueling DQN
- *Distributional RL*  
*(next time)*

## R2D2 (2018)

- Double DQN
- Prioritized buffer
- Multi-step DQN
  - + LSTM
- + Massive parallelization

# Frontiers

Rainbow DQN  
(2017)

- Double DQN
- Prioritized buffer
- Multi-step DQN
- Noisy Nets
- Dueling DQN
- *Distributional RL*  
*(next time)*

R2D2  
(2018)

- Double DQN
- Prioritized buffer
- Multi-step DQN
  - + LSTM
- + Massive parallelization

Agent 57  
(2020)

- Double DQN
- Prioritized buffer
- Multi-step DQN
  - + LSTM
- Massive parallelization
  - + Retrace
  - + Intrinsic motivation
  - + Meta-controller for hyperparameters

## Literature

- Playing Atari with Deep Reinforcement Learning (2013);
- *For DQN modifications use links on previous slide;*

